

Paper → Bibliography OS

The commands you actually run. Everything else is plumbing these call.

1 I already have everything

When: you have a folder (or tree) of papers — a finished literature review — and just want the bibliography.

Run: `/generate-bibliography <folder>`

Get: `bibliography.pdf` — one merged, deduplicated, clickable-DOI references PDF. One shot, runs in the background.

(This is the workflow that cleanly produced 86/87 correct references on the real 87-paper review.)

2 Talk in batches

When: your past papers are stored on disk, and you want to talk through new papers with Claude in batches (analysis, discovery) before producing the bibliography.

Run, repeatedly through the session: `/identify-papers-ai-ur <the papers you name>`

Each call resolves only the papers you reference and accumulates them (the AI cumulative).

Then, once: `/generate-bibliography-cumulative <folder>`

Get: `bibliography.pdf` — your on-disk folder *plus* everything you discovered in the batches, merged + deduplicated.

3 Walk-and-talk a section (disk-aware writing)

When: a bibliography already exists on disk, and you go on a walk voice-noting a draft over Telegram, citing papers as you go (some already on disk, some new).

Run: `/voice-writing-sample-all`

Cumulative sister: `/voice-writing-sample-all-cumulative` (draw references from the across-runs pool).

Get: an emailed writing-sample PDF. Reuses papers already on disk, resolves only the new ones, and numbers references [1], [2], [3]... in order of first citation.

Under the hood. Each command above is a *composite* — it fans out into ~20 other skills you never invoke directly (the `identify-papers-*`, `verify/validate`, `*-bibliography`, and `merge-*` families). They are the engine; these four are the dashboard. If you forget the rest, that's fine — you are meant to.

The four load-bearing commands:

`identify-papers-ai-ur`

`generate-bibliography-cumulative`

`generate-bibliography`

`voice-writing-sample-all`